

2021 - All about



Sanitizers, fuzzing and string-vectors

By Miklos Vajna

Software Engineer



Collabora
Online

vmiklos@collabora.com

About Miklos

From Hungary

- More details:
<https://www.collaboraoffice.com/about-us/>

Google Summer of Code 2010 / 2011

- Rewrite of the Writer RTF import/export

Then a full-time LibreOffice developer for SUSE

Now a contractor at Collabora



2021 - All about



Sanitizers

ubsan, asan and others

Clang provides several sanitizers, we use two:

- UndefinedBehaviorSanitizer (detects e.g. signed integer overflow)
- AddressSanitizer (detects e.g. stack-use-after-return and heap-use-after-free)

Environment

- core.git make check already passes with these sanitizers
- Now online.git make check (c++ tests) also pass
 - Cypress?
- Use LODE as the environment, as sanitizers have lots of config options, easy to hit non-interesting problems

2021 - All about



Fuzzing

Admin fuzzer

Tests the incoming websocket traffic of the admin console

- Simple file format: one websocket message / line
- Found 6 problems so far

```
Admin& admin = Admin::instance();  
auto handler = std::make_shared<AdminSocketHandler>(&admin);  
  
std::string input(reinterpret_cast<const char*>(data), size);  
std::stringstream ss(input);  
std::string line;  
while (std::getline(ss, line, '\n'))  
{  
    std::vector<char> v(line.data(), line.data() + line.size());  
    handler->handleMessage(v);  
}
```

Client session fuzzer

Initially this was “the fuzzer”, i.e. the first one:

- Tests what is incoming on the websocket from editing clients
- Found 11 problems so far

Fuzzer environment

- Same as sanitizers, i.e. ubsan+asan
- online.git configure gets an --enable-fuzzers
- Only uses Online as a library, i.e. the build produces no loolwsd binary
- The fuzzer is an executable, and it has to link all Online code statically

HTTP response fuzzer

Introduced as part of the async save work

- Tests what is a reply for a HTTP request
- Found 3 problems so far

Fuzzing-as-a-service

- All 3 fuzzers run 7/24 as a Jenkins job
- They run for a week: if they don't find anything, then they quit
 - Then pull, build, and start again
- Mail notification when they find something:
 - The server creates a reproducer (expensive)
 - A local environment can reproduce the produced crash sample (cheap)

String-vectors

Fuzzing found a pattern:

- If we have a vector of strings, it's easy to forget checking the array bounds before accessing the *n*th string
- If we are at it: allocating a null-terminated string for each token shows up on profiles

Solution: StringVector

- Similar to `std::vector<std::string>`, but it has a single underlying string
- Tokens only have offset + length “pointers” into that
- Safe API: if we would read past the end of the array, return an empty string
- Clang AST matcher to find all uses of `v[0] == “foo”`

Summary

Sanitizers: to make sure tests don't only pass by accident

- Have a tinderbox for this

Then fuzz it:

- Invent fake file formats to stress-test API that handles untrusted user input
- Do it as a CI job, so it finds badness before others do
- When the crash samples show a pattern, introduce safe APIs around unsafe ones

This makes Online a safer choice for everyone!